

Elasticsearch Guild

Analyzers & Edge N-gram Token Filter

Halil Bülent Orhon

Agenda

1. What is an inverted index?
2. What is an analyzer?
3. Built-in and custom analyzers
4. What is the Wildcard query?
5. What is the Edge N-gram tokenizer
6. Wildcard vs Edge N-gram

What is an Inverted Index?

- Maps **terms** to **document IDs**
- Enables fast full-text search

```
Term      → DocIDs
-----
cikolatali → 1, 3, 8
ve         → 1, 2, 5
fistikli   → 1
kek        → 1, 7
```

What is an Analyzer?

- Analyzer = `char_filter` → `tokenizer` → `token_filter`
- It transforms raw text into searchable tokens.
- Crucial for building the **inverted index**

Analyzer Usage: Index-time vs Query-time

- Analyzers are applied at **two points**:
 - **Index-time** → When documents are stored
 - **Query-time** → When user searches

 **Mismatch** between them can lead to no results!

Analyzer Component Summary

Component	Responsibility
<code>char_filter</code>	Pre-tokenization text normalization (e.g. replace <code>&</code> with <code>and</code>)
<code>tokenizer</code>	Splits text into tokens (e.g. words)
<code>token_filter</code>	Transforms tokens (e.g. lowercase, stemming, n-gram)

🧠 These three components form a complete **analyzer**

→ Customized analyzers = better search quality and control

Analyzer Pipeline – Step-by-Step Example

Input: "Çikolatalı & fıstıklı kek"

1. Char Filter

"&" → "ve"

→ "Çikolatalı ve fıstıklı kek"

2. Tokenizer (standard)

→ ["Çikolatalı", "ve", "fıstıklı", "kek"]

3. Token Filters (lowercase , asciifolding)

→ ["cikolatali", "ve", "fistikli", "kek"]

✅ Final tokens are stored in the inverted index

Some Built-in Analyzers

Analyzer	Description / Best Use Case
standard	General-purpose full-text search (default)
simple	Clean, letter-only content
whitespace	Tokenize without affecting punctuation
keyword	Exact-match fields
pattern	Custom delimiter-based tokenization

What is the **stop** Token Filter?

- Removes common words like **"and"**, **"or"**, **"the"**
- Helps reduce index size and noise

Example:

Input: **"çikolatalı ve fıstıklı kek"**

Output after **stop**: **["çikolatalı", "fıstıklı", "kek"]**

 Removing stopwords may break **phrase queries**

What is `word_delimiter` Token Filter?

Input: `"blue-jeans123"`

Output: `"blue"`, `"jeans"`, `"123"`, `"bluejeans123"`, `"blue-jeans123"`

```
"wordDelimiterTokenFilter": {  
  "type": "word_delimiter",  
  "catenate_all": true,  
  "split_on_numerics": true,  
  "preserve_original": true  
}
```

Creating a Custom Analyzer

Combine tokenizer, char filters, and token filters tailored to your language and use case.

You can define it under index settings:

→ see next slide for example

Example: Simplified Turkish Analyzer

```
"analyzer": {  
  "turkish_suggestion_analyzer": {  
    "type": "custom",  
    "tokenizer": "standard",  
    "char_filter": ["turkishCharFilter"],  
    "filter": [  
      "lowercase",  
      "asciifolding",  
      "word_delimiter"  
    ]  
  }  
},  
"char_filter": {  
  "turkishCharFilter": {  
    "type": "mapping",  
    "mappings": ["ı => i", "ş => s", "ç => c"]  
  }  
}
```

Wildcard Query

```
{
  "simple_query_string": {
    "fields": ["text"],
    "query": "elb*",
    "analyze_wildcard": true
  }
}
```

⚠ Requires scanning **all terms** → Slow

Wildcard Query - Performance Issues

- Query-time operation
- High CPU usage
- Not scalable
- Grows with vocabulary size

Edge N-Gram Analyzer – What It Does

- Indexes all prefixes of a word

Input: "elbise"

Indexed tokens: ["e", "el", "elb", "elbi", "elbis", "elbise"]

→ Matched directly at query-time

Why Use Edge N-Gram?

- Shifts cost to **index-time**
- Low latency at query-time
- Ideal for:
 - Search suggestions
 - Prefix search

Edge N-Gram Analyzer – Definition

```
"filter": {  
  "edgeNgramFilter": {  
    "type": "edge_ngram",  
    "min_gram": 1,  
    "max_gram": 20  
  }  
},  
"analyzer": {  
  "suggestionAnalyzer": {  
    "type": "custom",  
    "tokenizer": "standard",  
    "filter": [  
      "turkish_lowercase",  
      "wordDelimiterTokenFilter",  
      "asciifolding",  
      "edgeNgramFilter"  
    ]  
  }  
}
```

Search Analyzer

```
"analyzer": {  
  "suggestionSearchAnalyzer": {  
    "type": "custom",  
    "tokenizer": "standard",  
    "filter": [  
      "turkish_lowercase",  
      "wordDelimiterTokenFilter",  
      "asciifolding"  
    ]  
  }  
}
```

Mapping with Index + Search Analyzer

```
"mappings": {  
  "properties": {  
    "text": {  
      "type": "text",  
      "analyzer": "suggestionAnalyzer",  
      "search_analyzer": "suggestionSearchAnalyzer"  
    }  
  }  
}
```

Search Without Wildcard

```
{  
  "simple_query_string": {  
    "fields": ["text"],  
    "query": "elb"  
  }  
}
```

- ✓ "elb" is already indexed
- ✓ No wildcard needed

Performance Results

Metric	Wildcard	Edge N-Gram	Gain
Response Time	40.3ms	15.8ms	60% faster
Search Rate	1.4k / sec	63.6k / sec	44× higher
CPU Usage (peak)	82%	31%	62% lower
Index Size	6.8 GB	7.2 GB	+5.8%

Key Takeaways

- Wildcard is slow, not scalable
- Edge N-Gram shifts work to index-time
- Huge performance benefits with tiny index cost

Conclusion

✓ Edge N-Gram Analyzer provides:

- ⚡ Faster search latency
- 🧠 Lower CPU usage
- 📈 Higher throughput
- 💰 Reduced infrastructure needs

Recommended for:

- Search suggestions
- Autocomplete
- High-traffic search boxes

Thank you for listening!